



Qwen3 MemAgent RL / fixed external64 / batch-invariant TRITON_ATTN



Same source questions and decoding budgets; each model uses its own base. One training run per model.

Repeats are shown separately and do not increase the number of independent questions or training seeds.