



Qwen3-32B / six-GPU separate_async / 35/128 completed steps

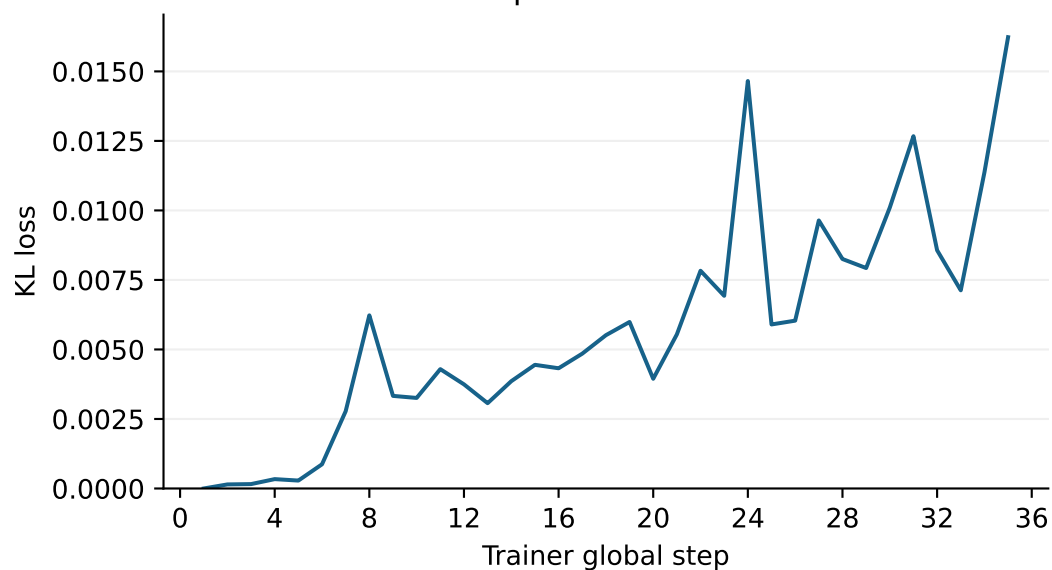
Training and waiting for the next batch



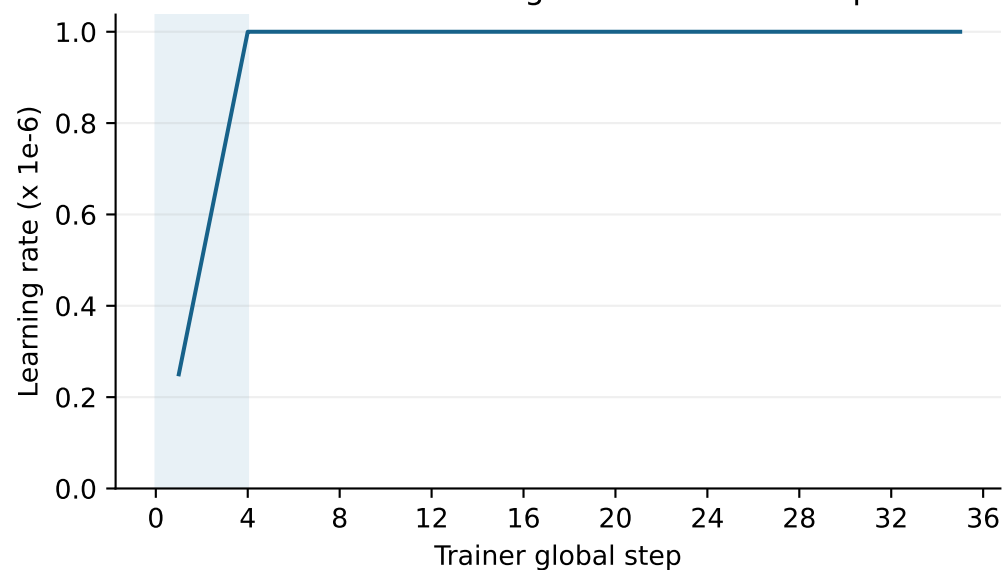
Gradient norm before clipping



Reported KL loss



Scheduler learning rate after each step



Batch wait is trainer timing_s/gen, not total serving time. LR shown is after the scheduler step.